

Solving Problems by Searching

CS 3600

Intro to Artificial Intelligence

Simplifying the domain

Our first intelligent agent will be restricted to work on environments that are

- **Known**
- **Fully observable**
- **Single-agent**
- **Deterministic**
- **Episodic***
- **Static**
- **Discrete**

Extensions that relax these assumptions

- **Unknown** (Learning agents)
- **Partially observable** (POMDPs)*
- **Multi-agent** (min-max search)
- **Stochastic** (probabilistic reasoning)
- **Sequential** (Hierarchical Planning)*
- **Dynamic** (Replanning)
- **Continuous** (RRTs and Controls)*

Search-based agent

```
def __init__(self, init_state):
    self.state = init_state
    self.problem = None
    self.plan = list()
    self.goal = None
def search_based_agent(self, percept):
    self.state = self.update_state(percept)
    if len(self.plan) <= 0:
        self.goal = self.make_goal(self.state)
        self.problem = self.make_problem(self.state, self.goal)
        self.plan = self.search(self.problem)
    action = self.plan.pop(0)
    return action
```

This agent is **offline**: it decides on a full plan of action before taking a single step

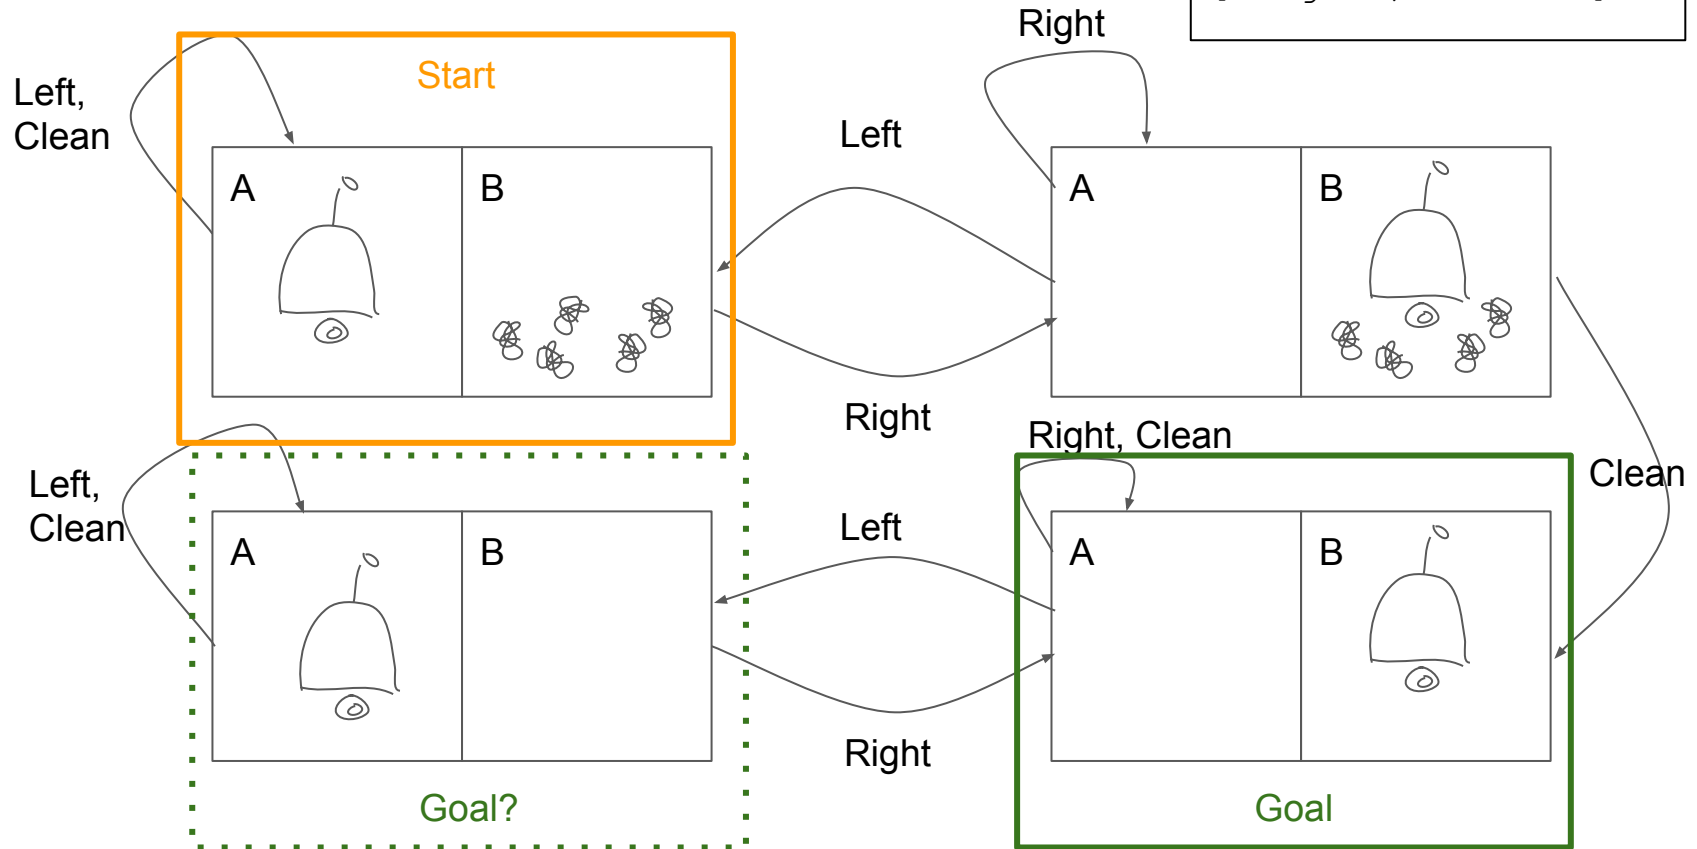
Components of a Search-based agent

- **update_state(percept)** - Construct a state representation from a given percept.

```
{ 'loc': 'A', status: 'clean' } →  
{ 'loc': 'A', 'A-clean': True, 'B-clean': self.state['B-clean'] }
```
- **choose_goal(state)** - Define success (goal state? goal test?)
make_problem(state, goal) - set up actions, construct state space (explicit? successor function?), initialize book-keeping
- **search(problem)** - returns a sequence of actions that take the agent from the start state to the/a goal state

Example problem: Vacuum World

Solution:
['Right', 'Clean']



Components of a problem

For non-trivial problems, we need a way to **generate** the state space without explicitly representing every node/edge

- **Start state:** The initial state, where the agent starts
- **Successor function:** $S(\text{state})$ returns a set of (action, successor state) pairs
- **Goal test function:** $\text{Goal}(\text{state})$ returns true if the state is a goal
- **Step cost:** $c(s_1, a, s_2)$ returns the cost of moving from s_1 to s_2 using action a


Components of a problem example: sudoku

- **Start state:**
Partially filled board
- **Successor function:**
 $S(\text{state})$: states generated by filling in one blank space with a non-conflicting number 1-9
- **Goal test function:**
 $\text{Goal}(\text{state})$: Is the entire board filled with non-conflicting numbers?
- **Step cost:**
 $c(s_1, a, s_2)$: 1

					1		5	
7	2		3	5			9	1
			8		7			
		8			5			4
	4	1		3		6	8	
5			4			7		
			6		3			
4	8			7	9		1	6
	9		1					

Components of a problem example: Romania

- **Start state:**
'Arad'
- **Successor function:**
 $S(\text{state})$: Neighboring cities of state
- **Goal test function:**
 $\text{Goal}(\text{state})$: $\text{state} == \text{'Bucharest'}$
- **Step cost:**
 $c(s_1, a, s_2)$: distance from s_1 to s_2 via 'a'




Representing search space - tree version


- For some problems, the number of states is too large (infinite?) to construct an explicit graph
- We can build the pieces of the state space we need to search 'as we go'
- The **search tree** is rooted at the initial state, leaves are expanded into their successors, may contain duplicate states (but not **nodes**!)
- Implementation note: children have 'back-pointers' to parents

We know how to search trees!


(a) The initial state




(b) After expanding Arad




(c) After expanding Sibiu



Depth First Search



Breadth First Search



Iterative Deepening Search


Limit = 0




Limit = 1





Limit = 2



Limit = 3



Uniform Cost Search (Dijkstra's)



Comparing search algorithms

	BFS	UCS	DFS	IDS
Time	$O(b^d)$	$O(b^{d+1}), O(b^{1+C^*/\epsilon})$	$O(b^m)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{d+1}), O(b^{1+C^*/\epsilon})$	$O(b^*m)$	$O(b^*d)$

b: branching factor, **d**: depth of shallowest solution

m: maximum depth of the tree, ϵ : smallest step cost, **C***: cost of optimal solution

Complete: BFS & IDS (if $b < \infty$), DFS (if $m < \infty$), UCS (if $\epsilon > 0$, and $b < \infty$)

Optimal: BFS & IDS (if all steps cost ϵ), UCS

Preview: Generic Search Algorithm

DFS, BFS, and UCS can be implemented with **a single algorithm!** Choice of data structure for the “next child to expand” determines which one.

- BFS: queue (children are expanded in the order they are added)
- DFS: stack (children are expanded in last-in-first-out order)
- UCS: priority queue (children are expanded based on cost-from-start)

IDS requires a small tweak: a depth limit parameter

Summary and preview

Wrapping up

- Search based agents work offline to find a sequence of actions that gets them from the initial state to a goal state
- A **search problem** can be represented explicitly as a graph, or implicitly by a start state, a successor function, a goal test function, and a cost function
- With this formulation, we can use any number of well known search algorithms to solve search problems

Preview

- Generic Search Algorithm, Worked Examples