


Heuristics

CS 3600

Intro to Artificial Intelligence

Romania - A* with straight line distance heuristic



Romania - A* with straight line distance heuristic

Arad (366)

[S(393),T(447),Z(449)]

[A]

Sibiu (393)

[R(413),F(415),T(447),Z(449),O(671)]

[A,S]

Rimnicu Vilcea (413)

[F(415),P(417),T(447),Z(449),C(526),O(671)]

[A,S,R]

Fagaras (415)


[P(417),T(447),Z(449),B(450),C(526),O(671)]

[A,S,R,F]

Pitesti (417)

[B(418),T(447),Z(449),B(450),C(526),C(625),O(671)]

[A,S,R,F,P]




Solution: [A->S, S->R, R->P, P->B]

Cost: 418

Why do we need consistency?

- Necessary in the proof to ensure that $f(n_i)$ never decreased
- Why might $f(n_i)$ decrease? Shortcuts!




Shortcuts - example (1)

Heuristic: $h(A)=15$, $h(B)=3$, $h(C)=10$, $h(D)=0$, $h(G)=0$

Admissible? ✓

Consistent?

$$\begin{aligned}h(n) &\leq c(n,a,n') + h(n') \\h(A) &\leq c(A, \rightarrow, B) + h(B) \\15 &\leq 10 + 3 \quad \text{✗}\end{aligned}$$



Shortcuts - example (2)


A(15)
[B(13), C(15)]
[A]

B(13)
[D(14), C(15)]
[A,B]

D(14)
[C(15), G(24)]
[A,B,D]

C(15)
[G(24)]
[A,B,D,C]

G(24)
[]
[A,B,D,C,G]



State	A	B	C	D	G
h(State)	15	3	10	0	0

Solution returned: [A->B,B->D,D->G], cost 24
Optimal solution: [A->C,C->D,D->G], cost 18

Fixing Generic Search to handle shortcuts (1)

What was the problem?

- When a shorter path to D was encountered, D was already in the “closed” list
- If we find a shorter path for a node in the closed list, we need to update its $g(n)$...
- And then all the $g(n)$ of the children of that node...
- Which may re-order the open priority queue...

Yuck

Fixing Generic Search to handle shortcuts (2)

Computationally

- **update parent of s** isn't so bad if we use backpointers
- **recompute g and resort open** has to trace back the new path!

One implementation

- Do DFS from start state, recompute g as you go
- Don't expand a node if it's already in open

Note for project 1: shouldn't be necessary as all the heuristics will be consistent or inadmissible anyway

```
Initialize 'current' node to start state
Initialize 'closed' as an empty list
Initialize 'open' as one of (stack, queue, priority queue)
while not( current['state'] is goal state):
    Add current['state'] to closed
    successors = successors of current['state']
    for s in successors:
        if not(s.state is in closed):
            Add new node for state to open
            elif s.cost+current['g'] < old cost to s:
                update parent of s
                recompute g and resort open
    current = next node in open that's not in closed
path = list()
while current has a parent:
    Add current['action'] to the front of path
    current = current['parent']
return path
```




Heuristic effectiveness

The **effective branching factor** (b^*) for a heuristic is a way of characterizing how **helpful** that heuristic is.

If A^* finds a solution at depth d expanding N nodes, then b^* is the branching factor that a uniform tree of depth d would need to contain $N+1$ nodes

$$N+1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

Lower effective branching factor indicates the heuristic will be effective in solving larger problems with reasonable computation time



Simple problem domain: 8-puzzle

Initial state: scrambled board

Goal state: tiles in numerical order

Actions: slide one tile into the blank spot (move the blank spot one tile)

Cost: 1 per move

Two heuristics (for comparison):

h_1 = # of misplaced tiles

h_2 = distance of all tiles to final position

7	2	4
5		6
8	3	1

Start State

$$\begin{aligned}h_1 &= 8 \\h_2 &= 26\end{aligned}$$

	1	2
3	4	5
6	7	8

Goal State

$$h_1 = h_2 = 0$$

Experiment - nodes expanded for 8-puzzle

Depth	2	4	6	8	10	12	14	16	18	20	22	24
IDS	10	112	680	6384	47127	3644035						
$A^*(h_1)$	6	13	20	39	93	227	539	1301	3056	7276	18094	39135
$A^*(h_2)$	6	12	18	25	39	73	113	211	363	676	1219	1641

- 100 random puzzles for each depth
- IDS didn't finish in time for $d > 12$
- Both h_1 and h_2 outperform IDS
- h_2 seems better than h_1 for $d > 6$
- Effective branching factor is relatively stable across problem sizes

Effective branching factor

IDS: 2.45 to 2.87

h_1 : 1.33 to 1.79

h_2 : 1.22 to 1.79

Is h_2 **always** better than h_1 ?

Yes!

For any node, $h_1(n) \leq h_2(n)$ (each out of place tile must move at least one space)

When comparing heuristics, if $h_a(n) \leq h_b(n)$ for all n , we say h_b **dominates** h_a

Since A^* with consistent heuristics will always expand every node with $f(n)=g(n)+h(n)<C^*$, we should try to make $h(n)$ as large as possible (still admissible and consistent, efficient to compute)

Admissible heuristics: $0 \leq h(n) \leq h^*(n)$


Heuristic design - problem relaxation

One way of generating heuristics is to use solutions to a version of the problem with **fewer constraints**

h_1 : path cost if tiles can “teleport” to the correct spot

h_2 : path cost if tiles can slide over one another

The relaxed problem has the same state space with additional edges: so the cost of a **solution** in the relaxed problem is **guaranteed** to be an **admissible** heuristic for the original problem



Heuristic design - composite heuristics

If we have a set of (admissible, consistent) heuristics that are non-dominated, we can combine them!

$$h(n) = \max\{ h_1(n), h_2(n), \dots, h_k(n) \}$$

Note that this new heuristic dominates all of the component heuristics

Heuristic design - Pattern databases

Idea: pre-compute the solution to a simpler sub-problem, and **store** the solution length. When searching the larger problem, match states against subproblem patterns and use the solution length as the estimate

8-puzzle example: solve the puzzle for a subset of the tiles. Different subsets yield different heuristics.

Since $\text{sum}(h_1, h_2, \dots) \geq \max(h_1, h_2, \dots)$, why don't we just add heuristics together?


*	2	4
*		*
*	3	1

Start State

	1	2
3	4	*
*	*	*

Goal State

Heuristic design - disjoint pattern databases



If you can split the problem into **disjoint** subproblems, where solving one does not **reduce** the cost of solving another, you can actually **add** the subproblem solution costs (instead of taking the max), but this can be trickier than you expect!

Summary and preview

Wrapping up

- We need consistency to ensure generic search expands in order of increasing $f(n)$. We can fix generic search to work even for inconsistent heuristics, but it can get messy.
- **Effective branching factor** is a useful way of quantifying and comparing heuristic “helpfulness”
- Several ways of designing heuristics: problem **relaxation**, **composite** heuristics, and **pattern databases**.

Next time

- Adversarial Search